



**Barcelona  
Supercomputing  
Center**  
*Centro Nacional de Supercomputación*



# **AutoParallel**

## **A Python module for automatic parallelization and distributed execution of affine loop nests**

8th Workshop on Python for High-Performance and Scientific Computing 2018

**Cristián Ramón-Cortés**

Ramon Amela

Jorge Ejarque

Philippe Clauss

Rosa M. Badia

# Motivation

## « Parallel Issues

- Identifying parallel regions
- Concurrency management
- Orchestrate execution

## « Distributed Issues

- Remote execution
- Data Transfers

**Going one step further to ease the development  
of distributed applications**

**So that any field expert can scale up an  
application to hundreds of cores**

# AutoParallel Annotation

“ A single Python decorator to parallelize and distributedly execute sequential code containing affine loop nests

Python  
decorator

```
from pycompss.api.parallel import parallel
```

```
@parallel()  
def matmul(a, b, c, m_size):  
    for i in range(m_size):  
        for j in range(m_size):  
            for k in range(m_size):  
                c[i][j] += np.dot(a[i][k], b[k][j])
```

Automatic  
taskification

NO data  
management

Sequential  
Code

NO resource  
management



Grid



Cluster



Cloud



# Outline

« Architecture

« Evaluation

« Loop taskification (Advanced feature)

« Conclusions and Future Work



# AutoParallel Architecture




**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación

# AutoParallel Annotation

## Taskification of affine loop nests at runtime

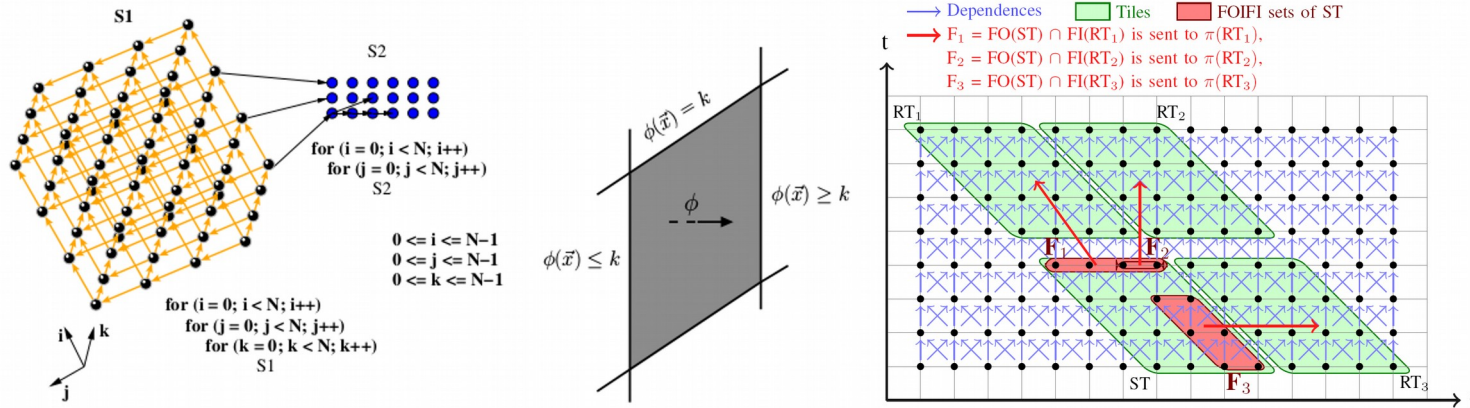
```
@parallel()  
def ep(mat, n, m, c1, c2):  
    for i in range(n):  
        for j in range(m):  
            mat[i][j] = compute(mat[i][j], c1, c2)
```



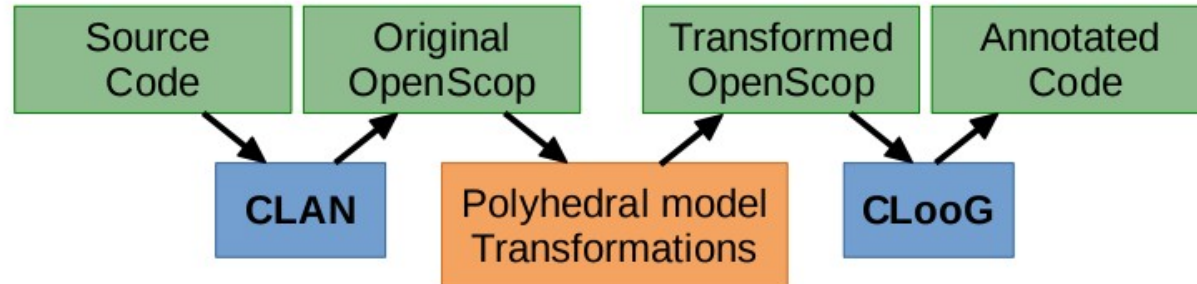
```
# [COMPSs AutoParallel] Begin Autogenerated code  
@task(var2=IN, c1=IN, c2=IN, returns=1)  
def S1(var2, c1, c2):  
    return compute(var2, c1, c2)  
  
def ep(mat, n, m, c1, c2):  
    if m >= 1 and n >= 1:  
        lbp = 0  
        ubp = m - 1  
        for t1 in range(lbp, ubp + 1):  
            lbv = 0  
            ubv = n - 1  
            for t2 in range(lbv, ubv + 1):  
                mat[t2][t1] = S1(mat[t2][t1], c1, c2)  
            compss_barrier()  
# [COMPSs AutoParallel] End Autogenerated code
```

# PLUTO Main Features

“ The Polyhedral model represents the instances of the loop nests’ statements as integer points inside a polyhedron.



“ PLUTO is an automatic parallelization tool based on the Polyhedral model to optimize arbitrarily nested loop sequences with affine dependencies.



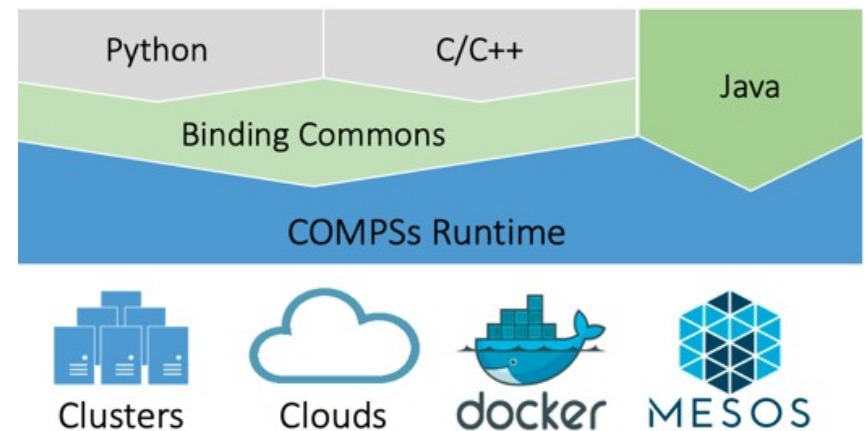
# PyCOMPSs Main Features

- COMPSs is a task-based programming model which aims to ease the development of parallel applications for distributed infrastructures
- The Python binding is known as PyCOMPSs



- Based on:
  - Sequential programming
  - Selection of tasks
    - Functions (instance and class methods)
    - Task data direction

```
@constraint(computingUnits="2")
@task(c=INOUT)
def multiply(a, b, c):
    c += a * b
```



- Same application runs on Clusters, Grids, Clouds and Containers



# AutoParallel Architecture

## Decorator:

- Implements the `@parallel()` decorator

## Python To OpenScop Translator:

- Builds a Python Scop object representing each affine loop nest detected in the user function

## Parallelizer:

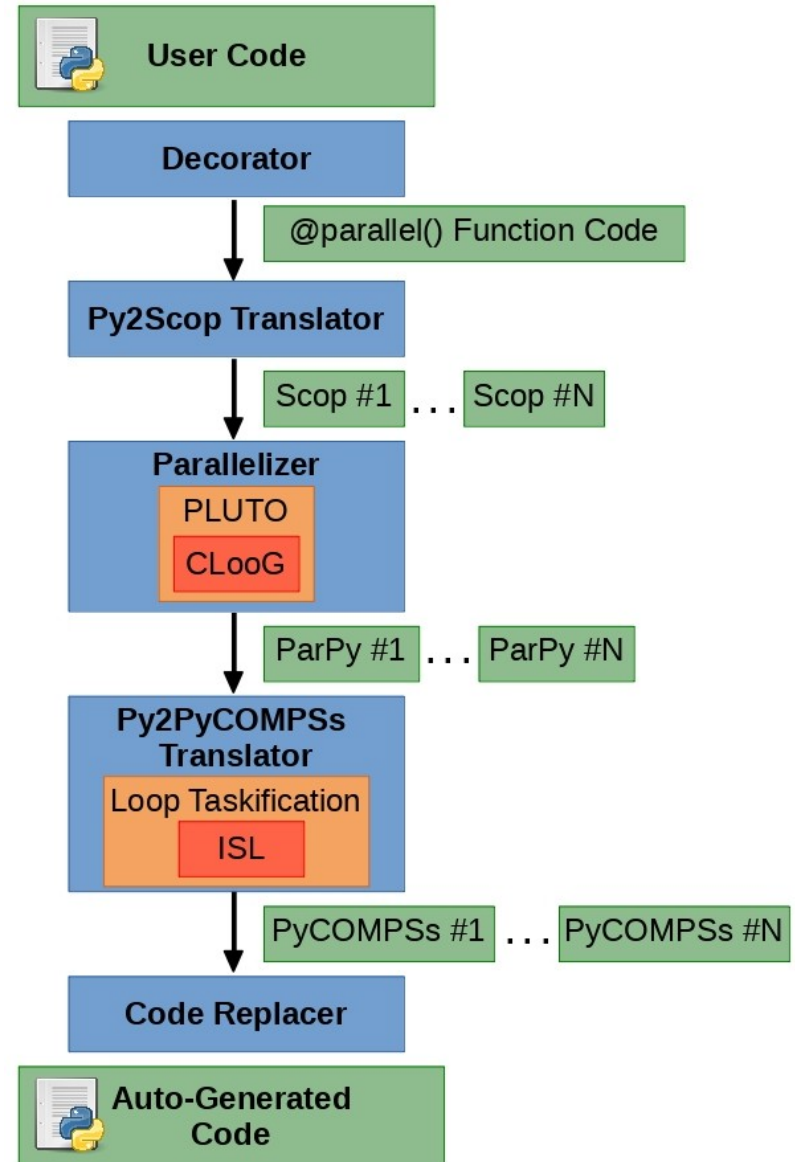
- Returns the Python code resulting from parallelizing an OpenScop file (OpenMP syntax)

## Python to PyCOMPSs Translator

- Inserts the PyCOMPSs syntax (task annotations and data synchronizations) to the annotated Python code

## Code Replacer

- Replaces each loop nest in the initial user code by the autogenerated code



# Evaluation



**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación


# Experimentation: Blocked Applications

## Cholesky

**LoC** Lines Of Code  
**CC** Cyclomatic Complexity  
**NPath** Npath Complexity

	Code Analysis		
	<i>LoC</i>	<i>CC</i>	<i>NPath</i>
<b>User</b>	220	26	112
<b>Auto</b>	274	36	14.576

	Loop Analysis		
	<i>#Main</i>	<i>#Total</i>	<i>Depth</i>
<b>User</b>	1	4	3
<b>Auto</b>	3	9	3



	Problem Size			Execution		
	<i>Total Matrix Size</i>	<i>#Blocks</i>	<i>Block Size</i>	<i>Task Types</i>	<i>#Tasks</i>	<i>SpeedUp @ 192 cores</i>
<b>User</b>	65.536 x 65.536	32 x 32	2048 x 2048	3	6.512	1,95
<b>Auto</b>				4	7.008	2,04


# Experimentation: Blocked Applications

LU

LoC Lines Of Code  
CC Cyclomatic Complexity  
NPath Npath Complexity

	Code Analysis		
	LoC	CC	NPath
User	238	35	79.872
Auto	320	39	331.776

	Loop Analysis		
	#Main	#Total	Depth
User	2	6	3
Auto	2	6	3



	Problem Size			Execution		
	Total Matrix Size	#Blocks	Block Size	Task Types	#Tasks	SpeedUp @ 192 cores
User	49.152 x 49.152	24 x 24	2048 x 2048	4	14.676	2,45
Auto				12	15.227	2,13




# Experimentation: Blocked Applications


## LU: In-depth Performance Analysis

- Paraver Trace with 4 workers (192 cores)

UserParallel



AutoParallel




# Experimentation: Blocked Applications

QR

LoC Lines Of Code  
CC Cyclomatic Complexity  
NPath Npath Complexity

	Code Analysis		
	LoC	CC	NPath
User	303	41	168
Auto	406	43	344

	Loop Analysis		
	#Main	#Total	Depth
User	1	6	3
Auto	2	7	3



	Problem Size			Execution		
	Total Matrix Size	#Blocks	Block Size	Task Types	#Tasks	SpeedUp @ 192 cores
User	32.768 x 32.768	16 x 16	2048 x 2048	4	19.984	2,37
Auto				20	26.304	2,10

# Loop Taskification: AutoParallel Advanced Feature



**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación

# PLUTO Tiling

- « Tiling a loop of given size  $N$  results in a division of the loop in  $N/T$  repeatable parts of size  $T$

## Original Loop

```
for i in range(N):  
    print(i)
```

## Tiled Loop

```
for i in range(N/T):  
    for i in range(T):  
        print(i*T + t)
```

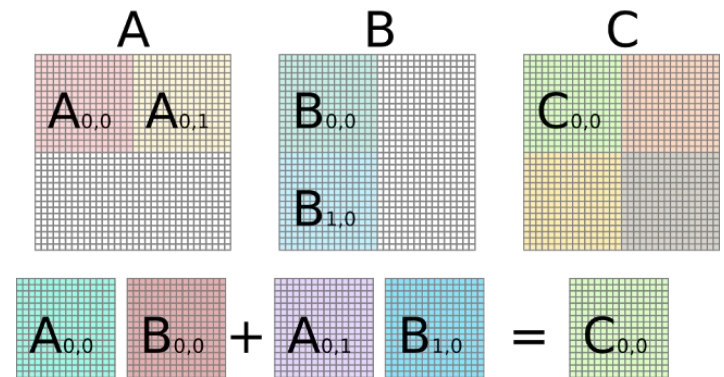
- « Tiling is designed to fit loops into the L1 or L2 caches  
**BUT** can be used to build data blocks to increase the tasks' granularity



# AutoParallel Loop Taskification

- Convert into tasks all the loops of a certain depth
- N-dimensional arrays are divided into data blocks (chunks) for each callee and reverted after the task execution

```
@parallel(pluto_extra_flags=["--tile"],
          taskify_loop_level=3)
def matmul(a, b, c, m_size):
    for i in range(m_size):
        for j in range(m_size):
            for k in range(m_size):
                c[i][j] += np.dot(a[i][k], b[k][j])
```



```
@parallel(pluto_extra_flags=["--tile"],
          taskify_loop_level=3)
def matmul(a, b, c, m_size):
    for i in range(m_size/T1):
        for j in range(m_size/T2):
            for k in range(m_size/T3):
                c[...][...] = LT1(c[...][...],
                                   a[...][...],
                                   b[...][...])
```

```
@task(..)
def LT1(c, a, b):
    for i' in range(T1):
        for j' in range(T2):
            for k' in range(T3):
                c[i'][j'] += np.dot(a[i'][k'],
                                     b[k'][j'])
    return c
```

# AutoParallel Loop Taskification

## EP Generated code

- Flattening and rebuilding data chunks

```
def ep(mat, n, m, c1, c2):
    if m >= 1 and n >= 1:
        lbp = 0
        ubp = m - 1
        for t1 in range(lbp, ubp + 1):
            lbv = 0
            ubv = n - 1
            # Chunk creation and flattening
            LT2_aux0 = [mat[t2][t1] for ...]
            LT2_au = ArgUtils()
            global LT2_args_size
            LT2_flat, LT2_args_size = LT2_au.flatten(LT2_aux0)
            # Task call
            LT2_ret = LT2(lbv, ubv, c1, c2, *LT2_flat)
            # Rebuild and re-assign
            LT2_aux_0, = LT2_au.rebuild(LT2_ret)
            ...
    compss_barrier()
```

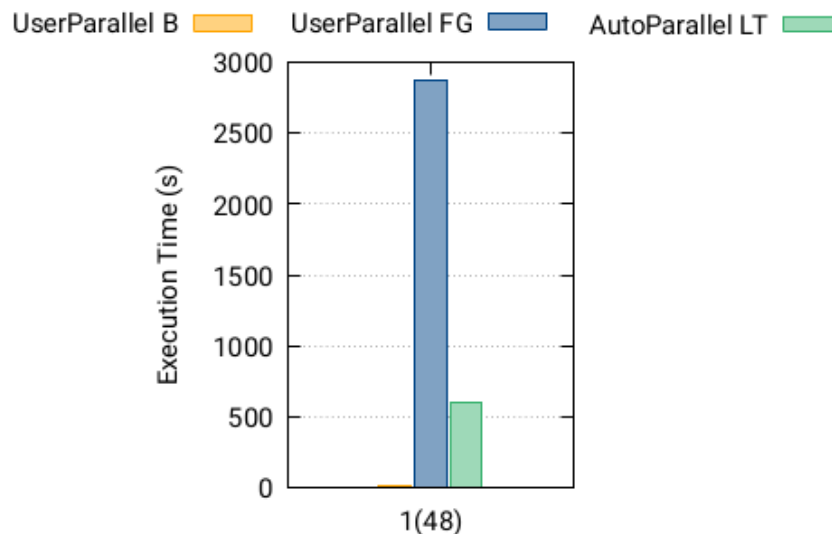
```
@task(lbv=IN, ubv=IN, c1=IN, c2=IN,
      returns="LT2_args_size")
def LT2(lbv, ubv, c1, c2, *args):
    global LT2_args_size
    var1, = ArgUtils.rebuild_args(args)
    for t2 in range(0, ubv + 1 - lbv):
        var1[t2] = S1_no_task(var1[t2],
                               c1, c2)
    return ArgUtils.flatten_args(var1)

def S1_no_task(var2, c1, c2):
    return compute(var2, c1, c2)
```

# Experimentation: Fine-grain Applications

## GEMM

	Code Analysis			Loop Analysis			Task Types
	<i>LoC</i>	<i>CC</i>	<i>NPath</i>	<i>#Main</i>	<i>#Total</i>	<i>Depth</i>	
User B	189	22	112	2	5	3	2
User FG	194	22	112	1	4	3	2
Auto LT	382	133	360064	2	4	3	4



# Conclusions and Future Work



**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación



# Conclusions and Future Work

## “ AutoParallel goes one step further in easing the development of distributed applications

- It is a Python module to automatically parallelize affine loop nests and execute them in distributed infrastructures
- The evaluation shows that the automatically generated codes for the Cholesky, LU, and QR applications can achieve the same performance than the manually parallelized versions

## “ Next steps

- Loop Taskification provides an automatic way to create blocks from sequential applications, but its performance is still far from acceptable.
  - Research on how to simplify the chunk accesses from the AutoParallel module.
  - Extend PyCOMPSs to support collection objects (e.g., lists)
- AutoParallel could be integrated with different tools similar to PLUTO to support a larger scope of loop nests.



**Barcelona  
Supercomputing  
Center**

*Centro Nacional de Supercomputación*



**EXCELENCIA  
SEVERO  
OCHOA**

# Thank you

Live demos at BSC booth #2038

Tuesday 4:40 pm

Wednesday 2:00 pm

Thursday 11:20 am



[cristianrcv/pycompss-autoparallel](https://github.com/cristianrcv/pycompss-autoparallel)



<http://compss.bsc.es/>

[cristian.ramoncortes@bsc.es](mailto:cristian.ramoncortes@bsc.es)